


Hybrid GraphRAG Manufacturing — System Architecture

Multi-turn chat • LangGraph-optional orchestration • Hybrid retrieval (BM25 + FAISS + KG) • Optional cause-ranker • Critic-validated tiered LLMs




Legend

Solid arrow — runtime data flow
Dashed arrow — reads from / writes to persisted state

Diagnostic-Mode Flow — LangGraph Topology

USE_LANGGRAPH=true routes every diagnostic query through this StateGraph. The procedural orchestrator follows the same logical flow.



Node	Function	Model / Tool	Env vars	Latency p50	Cost p50
format	Intent + entity extraction; query expansion	regex → qwen2.5:3b fallback	CLASSIFY_MODEL	5–60 ms	~\$0 (local)
retrieve	BM25 + FAISS + KG traversal → RRF top-k	rank-bm25 + FAISS + NetworkX	TOP_K_RETRIEVAL=10 TOP_K_RERANK=5 RRF50–300 ms		\$0 (no LLM)
rank_causes (opt-...	Score top-N root causes from evidence + KG	CAUSE_RANK_MODEL qwen2.5:3b	USE_CAUSE_RANKING CAUSE_RANK_TOP_K=5100–500 ms		\$0 (Ollama)
generate	First-pass evidence-grounded answer	ANSWER_MODEL gpt-4o	OPENAI_API_KEY	1.5–3 s	~\$0.0050
critic	Grounding / completeness / safety check	CRITIC_MODEL qwen2.5:3b	—	50–250 ms	\$0 (Ollama)
retry	Regenerate using critic feedback	RETRY_MODEL gpt-4o	MAX_CRITIC_RETRIES=2	1.5–3 s	~\$0.0050 (per r...

Cost & Latency Breakdown

Per-query estimates with the default tiered routing (answer = gpt-4o, critic = qwen2.5:3b on Ollama). Local models are free; OpenAI pricing per core/llm_client.py.

1. Per-mode summary (typical query, default config)

Mode	LLM calls per query	Tokens (in / out)	Cost / query	Latency p50	Notes
Quick Search	0	0 / 0	\$0.0000	120–400 ms	FAISS + embeddings only; w...
Diagnostic (default)	1 answer + 1 critic	~1500 / ~400 + ~700 / ~150	~\$0.0080	2.0–4.5 s	ANSWER on gpt-4o, CRITIC o...
Diagnostic + cause-ranker	+1 cause-ranker	+~1200 / ~250	+ \$0.0000	+100–500 ms	USE_CAUSE_RANKING=true; in...
Diagnostic worst-case	1 answer + 2 critics + 1 ret...	~3700 / ~950	~\$0.0160	4.5–9 s	Critic FAILs once, retry r...
Chat (multi-turn)	Same as Diagnostic	Same as Diagnostic	Same as Diagnos...	Same as Diagnost...	Plus per-turn slot-filling...
Classical RAG (baseline)	1	~1500 / ~400	~\$0.0005	1.0–2.0 s	FAISS-only retrieval, gpt-...
Direct LLM (baseline)	1	~150 / ~400	~\$0.0003	0.8–1.5 s	No retrieval; gpt-4o-mini ...

2. Per-stage detail for Diagnostic mode (USE_CAUSE_RANKING=true, 1 retry path)

Stage	Model	Provider	Tokens (in / out)	Cost / call	Latency p50	Mandatory?
format	regex → qwen2.5:3b	Ollama (fallback)	0–250 / 0–80	\$0	5–60 ms	yes
retrieve	—	FAISS + BM25 + KG	0 / 0	\$0	50–300 ms	yes
rank_causes	qwen2.5:3b	Ollama	~1200 / ~250	\$0	100–500 ms	USE_CAUSE_RANKIN...
generate	gpt-4o	OpenAI	~1500 / ~400	~\$0.0050	1.5–3 s	yes
critic (1st)	qwen2.5:3b	Ollama	~700 / ~150	\$0	50–250 ms	yes
retry (if FAIL)	gpt-4o	OpenAI	~1700 / ~400	~\$0.0050	1.5–3 s	FAIL & attempts<...
critic (2nd)	qwen2.5:3b	Ollama	~700 / ~150	\$0	50–250 ms	after retry


3. Cloud vs local pricing (per 1k tokens, from core/llm_client.py)

Model role	Default	OpenAI in / out	Local equivalent	Switch via .env	Net effect
Answer / Retry	gpt-4o	\$0.0025 / \$0.010	qwen2.5:7b on Ollama	ANSWER_MODEL=qwen2.5:7b	100% free, slower late...
Critic	qwen2.5:3b	n/a (local)	—	CRITIC_MODEL=gpt-4o-mini	\$0 → ~\$0.0001 / call
Cause-ranker (opt-in)	qwen2.5:3b	n/a (local)	—	CAUSE_RANK_MODEL=gpt-4o-mini	\$0 → ~\$0.0002 / call
Comparison (Direct/Classical)	gpt-4o-mini	\$0.00015 / \$0.0006	qwen2.5:3b on Ollama	DIRECT_LLM_MODEL / CLASSICAL_RAG_MODEL	Free at slight quality...
Embeddings	MiniLM-L6-v2	n/a (local)	—	EMBEDDING_MODEL=BAAl/bge-small-en-v1.5	Rebuild FAISS after ch...

Human-in-the-Loop (HITL) Approval Gate

criticality_check + human_approval (interrupt) • SQLite checkpointer • audit log • USE_HITL=true

LangGraph topology with the HITL gate (Phases A + B + C)



Risk score drivers — core/criticality_classifier.py

Driver	Trigger	Score bump	Domain
safety_keyword:*	Substring of HITL_HIGH_RISK_KEYWORDS in query or proposed answer	0.55 + 0.05 / extra hit	diagnostic
high_risk_intent	Clarifier intent ∈ {shutdown, emergency, lockout_tagout, permit_to_work}	0.90	diagnostic
low_critic_confidence	Critic verdict.confidence < 0.5	0.30	diagnostic
purchase_value≥threshold	PurchaseRequest.total_usd ≥ HITL_AUTO_APPROVE_BELOW_USD	0.70	purchase_request
single_source_vendor	KG flags the part with single_source=true	0.65	purchase_request
long_lead_time	lead_time_days > 7	0.55	purchase_request
class_A_equipment	Used by an Equipment node tagged criticality=A	0.70	purchase_request
llm_grader:*	Tier-2 LLM grader (only fires for inconclusive 0.3–0.7 band)	max(score, llm_score)	any

REST surface (api/server.py)

GET /api/approvals/pending list paused HITL workflows
GET /api/approvals/{thread_id} snapshot of one paused workflow
POST /api/approvals/{thread_id}/resume {approved, approver, comments, edited_answer?}
GET /api/audit?limit=N&offset=M recent decisions + approval-rate stats

Pipeline status state machine

in_progress → awaiting_approval → complete (approved)
in_progress → awaiting_approval → rejected (approved=false)
in_progress → complete (auto-approve / USE_HITL)

Checkpoint: SqliteSaver (HITL_DB_PATH) — survives restarts.
Audit log: core/audit_log.py — append-only, one row per decision